

Overview

python3 train.py

319.6%

peak cpu

40.49s

cpu time

586.6MiB

peak memory

20.50s

duration

4

processes

mode

run

started

2026-09-14T08:00:12.764+00:00

interval

0.5s requested

samples

42

host

devbox · 32 cpus · 31.0GiB ram

boundary

treehawk-65397-1789372812.scope

peak concurrent

4

mean cpu

197.5%


peak rss / pss

606.0MiB / 586.2MiB

exit code


0

how each process was found




CPU over time

100% is one core fully used; the workload total includes processes that exited between samples





Memory over time

rss double-counts pages shared between children; pss divides them fairly; cgroup is the kernel's own figure




Process lifetimes

each bar is one process, from the first sample it appeared in to the last




CPU by process

stacked; a process contributes nothing in the sample it first appears in, because there is no earlier counter to compare





Memory by process

stacked rss; shared pages are counted once per process, so this reads high



Biggest consumers


totals over the whole run, not a single instant



Sampling quality

polling cannot see a process that starts and ends between two samples; these show how tight the sampling actually was

sampling error · worst +0.1 ms



cpu distribution

